

# 18.S096 January 2017: Memory and Matrices

Steven G. Johnson, MIT Applied Math

performance experiments (circa 2008):

Hardware: 2.66GHz Intel Core 2 Duo

64-bit mode, double precision, gcc 4.1.2

optimized BLAS dgemm: ATLAS 3.6.0

<http://math-atlas.sourceforge.net/>

# A trivial problem?

$$\underset{m \times p}{C} = \underset{m \times n}{A} \underset{n \times p}{B}$$

the “obvious” C code (rows · columns):

```
/* C = A B, where A is m x n, B is n x p,
   and C is m x p, in row-major order */
void matmul(const double *A, const double *B,
            double *C, int m, int n, int p)
{
    int i, j, k;
    for (i = 0; i < m; ++i)
        for (j = 0; j < p; ++j) {
            double sum = 0;
            for (k = 0; k < n; ++k)
                sum += A[i*n + k] * B[k*p + j];
            C[i*p + j] = sum;
        }
}
```

**for**  $i = 1$  **to**  $m$

**for**  $j = 1$  **to**  $p$

$$C_{ij} = \sum_{k=1}^n A_{ik} B_{kj}$$

**$2mnp$  flops**


(adds+mults)

“floating-point  
operations”

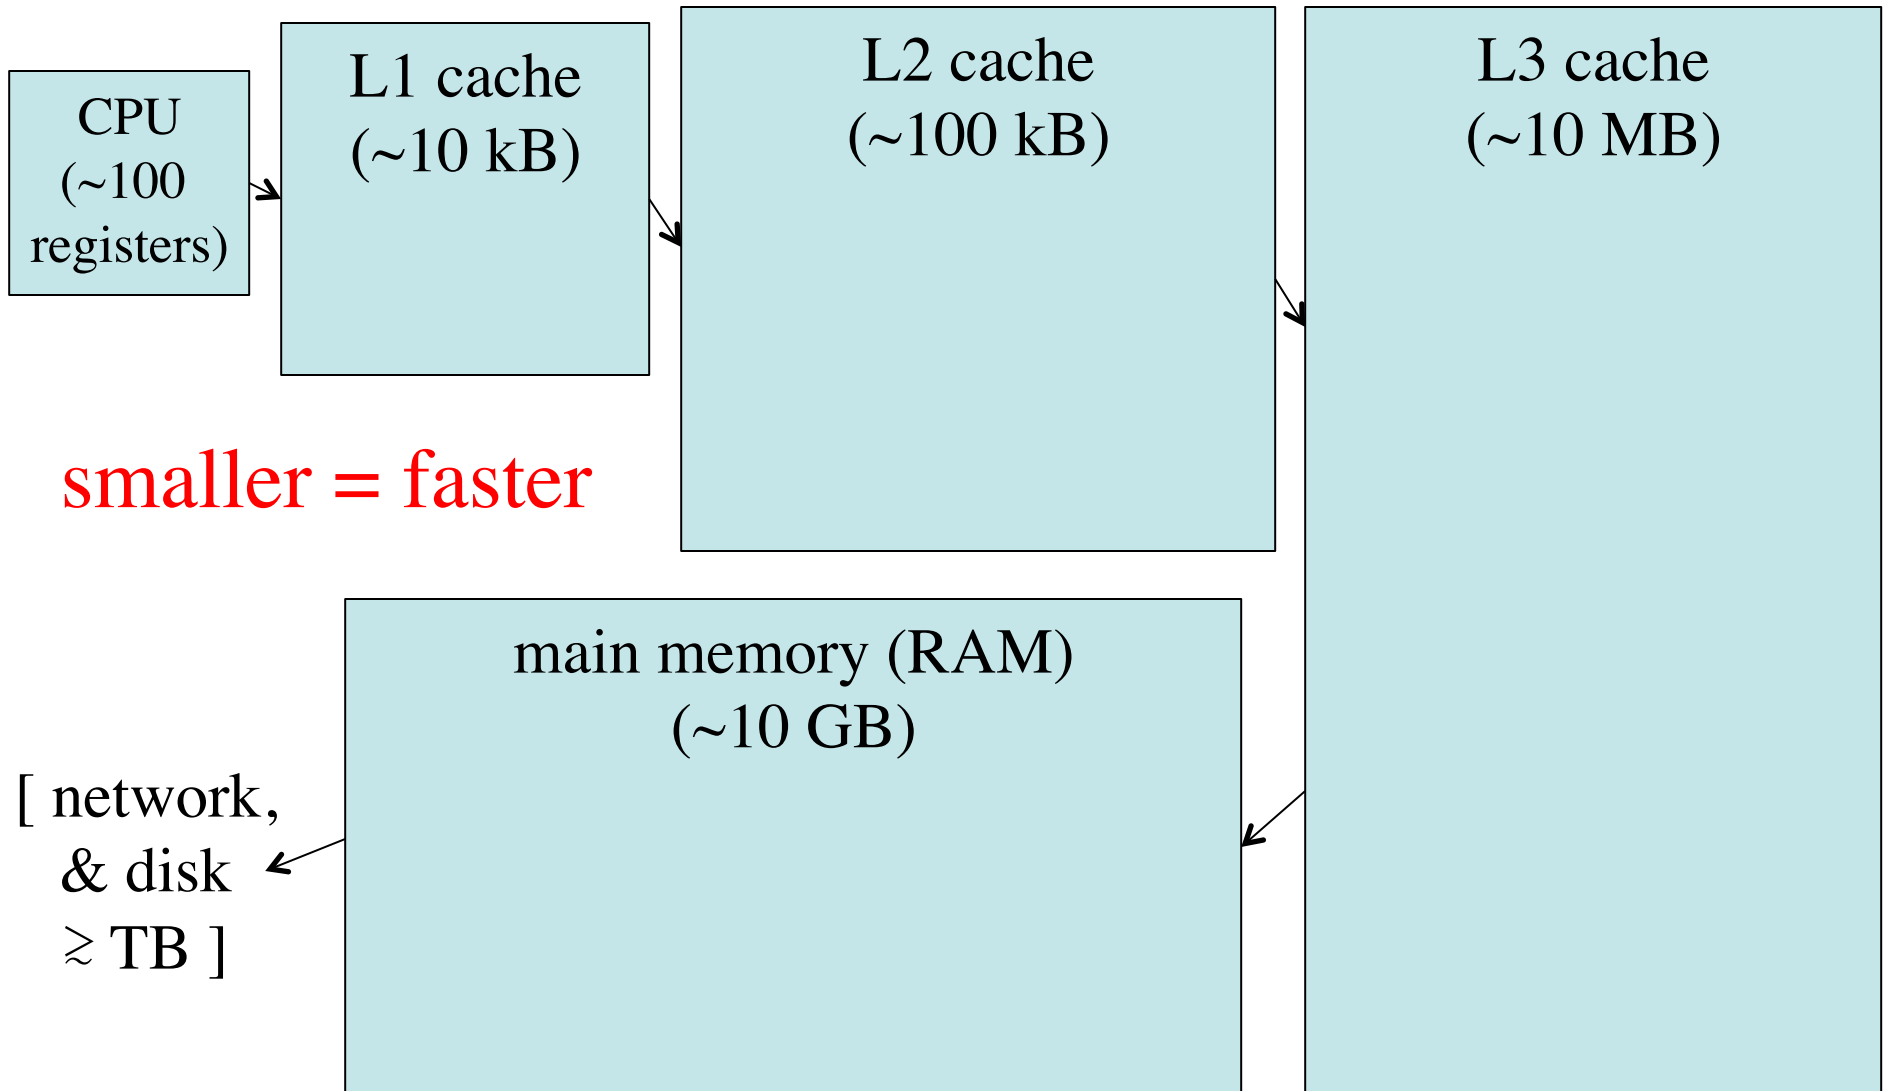
just three loops, how complicated can it get?

# flops/time is not constant!


(square matrices,  $m=n=p$ )



# Speed is limited by access to the **memory hierarchy** *[not to scale!]*



# All flops are not created equal




# Things to remember

- We often **cannot understand performance without understanding memory** efficiency (caches).
  - often  $\sim 10$  times more important than arithmetic count when working with lots of data
- Computers are **more complicated than you think**.
- Even a trivial algorithm is nontrivial to implement *well*.
  - matrix multiplication: 10 lines of code  $\rightarrow$  **130,000+** (ATLAS)
  - getting the **last factor of 2 in speed** often requires wizardry **(and is usually not worth it)**
  - but **factors of 10** are often worthwhile and not too hard...

# Ideal Cache Model

[ Frigo, Leiserson, Prokop, and  
Ramachandran (1999) ]

simplified model of cache to help us understand/design algorithms



when CPU needs an item, either:

- **cache hit: already in cache** (fast)
- **cache miss: load into cache** (slow)

goal: analyze **# of cache misses**

Simplification: **cache is “ideal”**

- **fully associative**: any item in memory can replace any item in cache
- **optimal replacement**: cache miss replaces “best” item  
(= item not needed for longest time in the future)  
... within a  $\sim$ constant factor of more realistic caches

**Cache complexity**: for problem with  $n$  items, cache size  $Z$ ,  
want **# misses for large  $n$**  in “big O” or “big  $\Theta$ ” notation  
(**ignoring constant factors**), e.g.  $\Theta(n/Z)$


*Strategy for efficient cache utilization:*  
**Maximize temporal locality**

Once we read an item from memory, we want as much computation as possible before reading the next item...

Re-arrange our algorithm so that computations on the same data occur at close to the same time.



# (optimal) Blocked Matrix Multiply



divide matrices into  $b \times b$  blocks of  $b = \sqrt{Z/3}$  numbers  
 load 3  $b \times b$  blocks into cache and *multiply blocks in-cache*


(provably optimal) cache misses:  $\Theta(n^3/b^3)$   $\times$   $\Theta(b^2)$  =  $\Theta(n^3/\sqrt{Z})$

# block  $\times$  block      # cache misses  
 multiplications          per block

# Challenges with blocking

- Programmer/code **needs to know the size  $Z$**  of the cache — different code for every CPU?
- Multiple levels of cache = **nested blocking**
- Many complications to get near-optimal “**constant factor**” in  $\Theta$ 
  - optimal block size is non-square to balance load/store cost.
  - lowest level (cache=registers) requires unrolling, SIMD optimizations, lots of tricks...

# (optimal) Cache-Oblivious Matrix Multiply



*divide and conquer:*

divide  $C$  into 4 blocks

compute block multiply recursively

achieves **optimal**  $\Theta(n^3/\sqrt{Z})$  cache complexity  
without knowing the  $Z$ , works for **nested caches** too

# A little C implementation (~25 lines)

```
/* C = C + AB, where A is m x n, B is n x p, and C is m x p, in
row-major order.  Actually, the physical size of A, B, and C
are m x fdA, n x fdB, and m x fdC, but only the first n/p/p
columns are used, respectively. */
void add_matmul_rec(const double *A, const double *B, double *C,
                   int m, int n, int p, int fdA, int fdB, int fdC)
{
    if (m+n+p <= 48) { /* <= 16x16 matrices "on average" */
        int i, j, k;
        for (i = 0; i < m; ++i)
            for (k = 0; k < p; ++k) {
                double sum = 0;
                for (j = 0; j < n; ++j)
                    sum += A[i*fdA + j] * B[j*fdB + k];
                C[i*fdC + k] += sum;
            }
    }
    else { /* divide and conquer */
        int m2 = m/2, n2 = n/2, p2 = p/2;

        add_matmul_rec(A, B, C, m2, n2, p2, fdA, fdB, fdC);
        add_matmul_rec(A+n2, B+n2*fdB, C, m2, n-n2, p2, fdA, fdB, fdC);

        add_matmul_rec(A, B+p2, C+p2, m2, n2, p-p2, fdA, fdB, fdC);
        add_matmul_rec(A+n2, B+p2+n2*fdB, C+p2, m2, n-n2, p-p2, fdA, fdB, fdC);

        add_matmul_rec(A+m2*fdA, B, C+m2*fdC, m-m2, n2, p2, fdA, fdB, fdC);
        add_matmul_rec(A+m2*fdA+n2, B+n2*fdB, C+m2*fdC, m-m2, n-n2, p2, fdA, fdB, fdC);

        add_matmul_rec(A+m2*fdA, B+p2, C+m2*fdC+p2, m-m2, n2, p-p2, fdA, fdB, fdC);
        add_matmul_rec(A+m2*fdA+n2, B+p2+n2*fdB, C+m2*fdC+p2, m-m2, n-n2, p-p2, fdA, fdB, fdC);
    }
}

void matmul_rec(const double *A, const double *B, double *C,
               int m, int n, int p)
{
    memset(C, 0, sizeof(double) * m*p);
    add_matmul_rec(A, B, C, m, n, p, n, p, p);
}
```

**note: base case is  $\sim 16 \times 16$**

*recurring down to  $1 \times 1$*

*would kill performance*


(1 function call per element,

no register re-use)


dividing  $C$  into 4

— note that, instead, for  
very non-square matrices,  
we might want to divide  
 $C$  in 2 along longest axis

# No Cache-based Performance Drops!



...but absolute performance still sucks



# Registers == Cache

- The registers ( $\sim 100$ ) form a very small, almost ideal cache
  - Three nested loops is not the right way to use this “cache” for the same reason as with other caches
- Need long blocks of unrolled code: load blocks of matrix into local variables (= registers), do matrix multiply, write results
  - Loop-free blocks = many optimized hard-coded base cases of recursion for different-sized blocks ... often automatically generated (ATLAS)
  - Unrolled  $n \times n$  multiply has  $(n^3)!$  possible code orderings — compiler cannot find optimal schedule (NP hard) — cache-oblivious scheduling can help (c.f. FFTW), but ultimately requires some experimentation/wizardry (automated in ATLAS)
  - Optimal blocks are non-square to balance load/store cost, and details (e.g. scheduling) turn out to depend on the CPU.

No data re-use = no possibility of temporal locality ... what then?

Suppose we are computing the dot product  $\mathbf{x}^* \mathbf{y}$  of two vectors:

$$\mathbf{x}^* \mathbf{y} = \sum_{i=1}^n \overline{x_i} y_i$$

Each element of  $\mathbf{x}$  and  $\mathbf{y}$  is **used exactly once**.

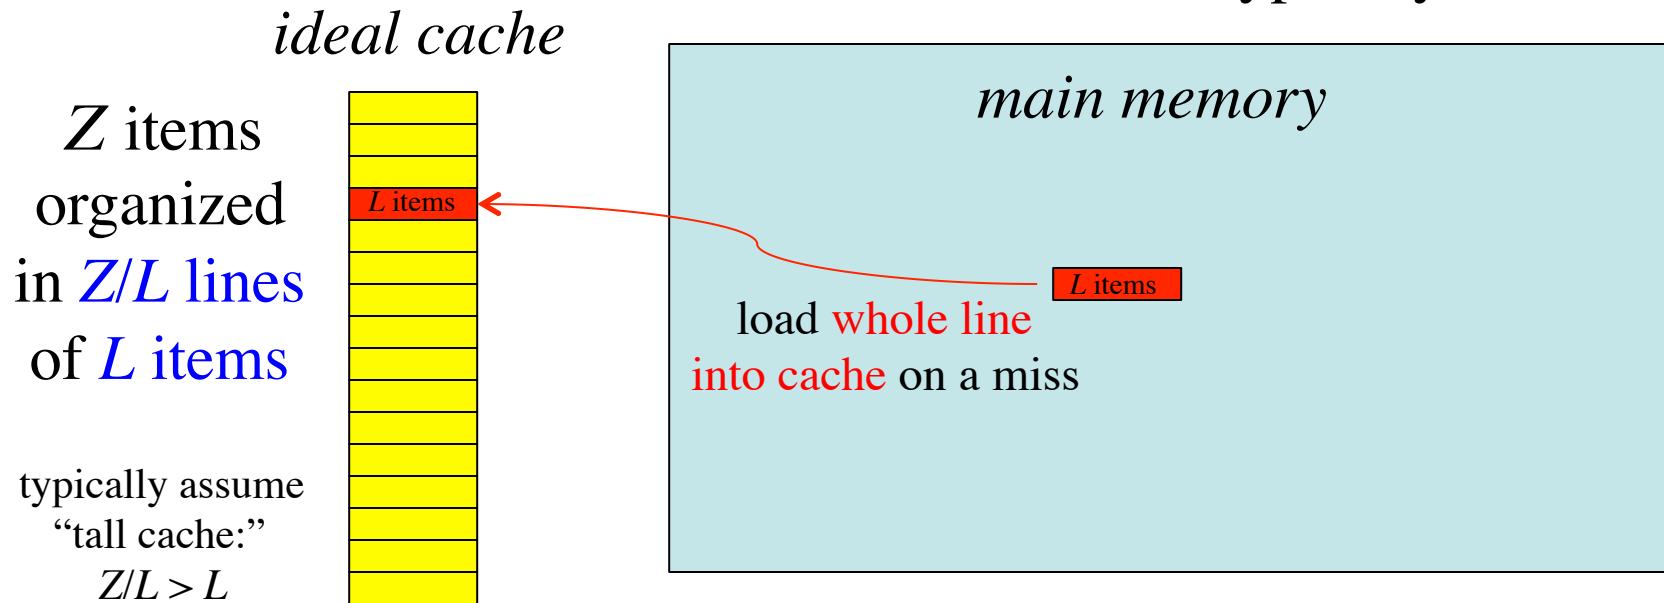
**Does this mean we get no benefit from the caches?**



# Cache lines and spatial locality

To speed up algorithms with little or no data re-use, caches exploit the fact that **memory access is often consecutive** by reading in a **whole cache line of  $L$  items on a miss**

[ typically,  $L \sim 64$  bytes ]



## Cache-optimization strategy:

when you access data in memory, **try to access *nearby data***  
soon afterwards ... maximize “**spatial locality**”

# Example: Matrix addition

$$\begin{array}{ccc} \boxed{\phantom{000000}} & = & \boxed{\phantom{000000}} + \boxed{\phantom{000000}} \\ C & & A \quad B \\ m \times n & & m \times n \quad m \times n \end{array}$$

two possible algorithms:

for  $i = 1$  to  $m$

for  $j = 1$  to  $n$

$$C_{ij} = A_{ij} + B_{ij}$$

or

for  $j = 1$  to  $n$

for  $i = 1$  to  $m$

$$C_{ij} = A_{ij} + B_{ij}$$

which one to use? depends on how matrices are stored!


# Column-major storage

used by Fortran, Matlab, Julia, ...

= store **columns** of  $A$  **consecutively** in memory

0	5	10
1	6	11
2	7	12
3	8	13
4	9	14

$A$   
 $5 \times 3$



access rows consecutively:

for  $i = 1$  to  $m$

for  $j = 1$  to  $n$

$$C_{ij} = A_{ij} + B_{ij}$$

=  $\Theta(mn)$  misses

access **columns** consecutively:

for  $j = 1$  to  $n$

for  $i = 1$  to  $m$

$$C_{ij} = A_{ij} + B_{ij}$$

=  $\Theta(mn/L)$  misses